

Restricted Boltzmann Machines

Theory and Implementation

Vagram Airlian
Software Engineer
LIT JINR
27.11.2020



Theory

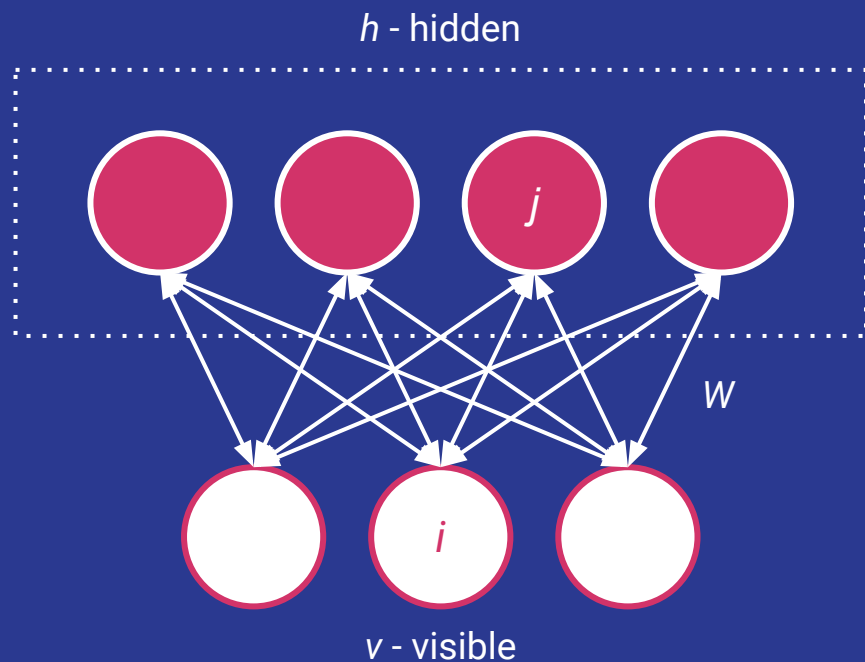
Restricted Boltzmann Machines

- An RBM defines a probability distribution over binary-valued patterns.
- They're undirected and don't have an output layer.
- All the hidden and visible nodes are all connected with each other.
- Restricted Boltzmann machines can generate data (a generative model).



RBM Model

- Not fully observable.
- One layer of hidden units (more is possible).
- No connections between hidden or visible units.
- A complete bipartite graph (biclique).



Gibbs Sampling


$$\Pr(x_i = 1 \mid x_{-i}) = \sigma \left(\sum_{j \neq i} w_{ij} x_j + b_i \right)$$

if $t \rightarrow \infty$, the configurations will be distributed approximately according to the model distribution.



Conditional Independence

- Provides an unbiased sample from the posterior distribution when given a data-vector.
- Allows easy vectorisation of computations.



First Step


$\mathbf{v}^{(0)}$ - initial (input) data

$$\mathbf{h}^{(0)} = \mathbb{E}[\mathbf{h} \mid \mathbf{v}] = \sigma(\mathbf{W}\mathbf{v}^{(0)} + \mathbf{b}_h)$$


$$\mathbf{v}^{(1)} = \mathbb{E}[\mathbf{v} \mid \mathbf{h}] = \sigma(\mathbf{W}^T \mathbf{h}^{(0)} + \mathbf{b}_v)$$

where

$\mathbf{v}^{(1)}$ - reconstructed data

$$\sigma(x) = (1 + e^{-x})^{-1}$$


First Step



Derivation

A joint configuration energy:

$$E(\mathbf{v}, \mathbf{h}) = -\sum_i \mathbf{a}_i v_i - \sum_i \mathbf{b}_i h_i - \sum_{i,j} v_i h_j w_{ij}$$

The network assigns a probability to every pair via:

$$p(\mathbf{v}, \mathbf{h}) = Z^{-1} e^{-E(\mathbf{v}, \mathbf{h})}, p(\mathbf{v}) = Z^{-1} \sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})} \text{ - for a visible vector, where } Z = \sum_{\mathbf{v}, \mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}$$

Then, the derivative of the log probability:

$$\partial \log p(\mathbf{v}) / \partial w_{ij} = \mathbb{E}_{\text{data}}[v_i h_j] - \mathbb{E}_{\text{model}}[v_i h_j]$$

and the learning rule:

$$\Delta w_{ij} = \varepsilon (\mathbb{E}_{\text{data}}[v_i h_j] - \mathbb{E}_{\text{model}}[v_i h_j]), \text{ where } \varepsilon \text{ - learning rate}$$

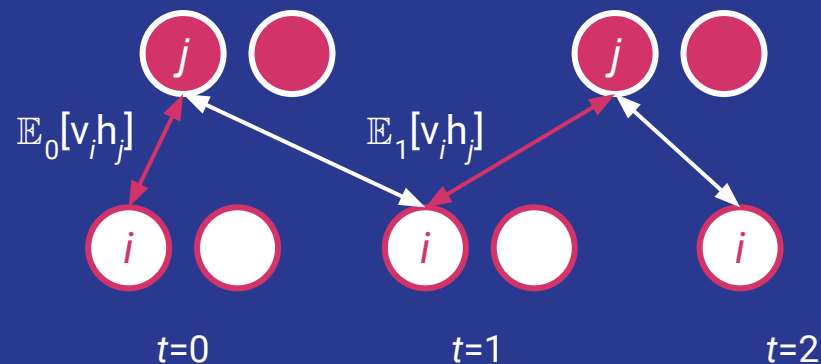


Contrastive Divergence

- Start with a training vector on the visible units.
- Update all the hidden units in parallel.
- Update the all the visible units in parallel to get a “reconstruction”.
- Update the hidden units again.
- Repeat k times.

$$\Delta w_{ij} = \varepsilon (\mathbb{E}_{\text{data}}[v_i h_j] - \mathbb{E}_{\text{reconstruction}}[v_i h_j])$$

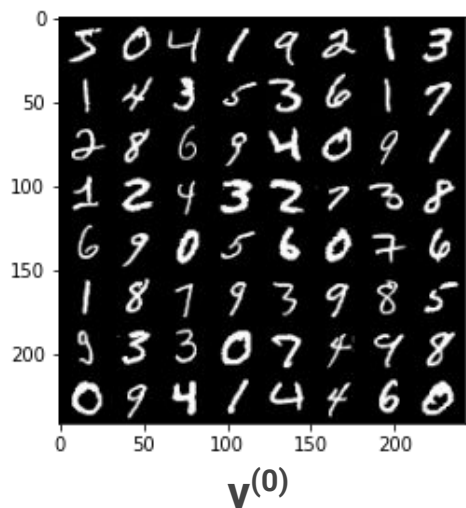
(because t is finite)




$$\mathbb{E}_{\text{data}}[v_i h_j] = \mathbb{E}_0[v_i h_j]$$
$$\mathbb{E}_{\text{reconstruction}}[v_i h_j] = \mathbb{E}_1[v_i h_j]$$

MNIST Example

Training Sample




Generated Sample, $k=2$



Sample Features

Samples of the learned features



Netflix Challenge

You are given most of the ratings that half a million Users gave to 18 000 movies on a scale from 1 to 5 (stars).

- Each user only rates a small fraction of the movies.
- You have to predict the ratings users gave to the U4 held out movies.

| Users | Movies (18 000) | | | | | |
|-------|-----------------|----|----|----|----|----|
| | M1 | M2 | M3 | M4 | M5 | M6 |
| U1 | | | | 3 | | |
| U2 | 5 | | 1 | | | |
| U3 | | 3 | 5 | | | |
| U4 | 4 | | ? | | | 5 |
| U5 | | | 4 | | | |
| U6 | | | | 2 | | |



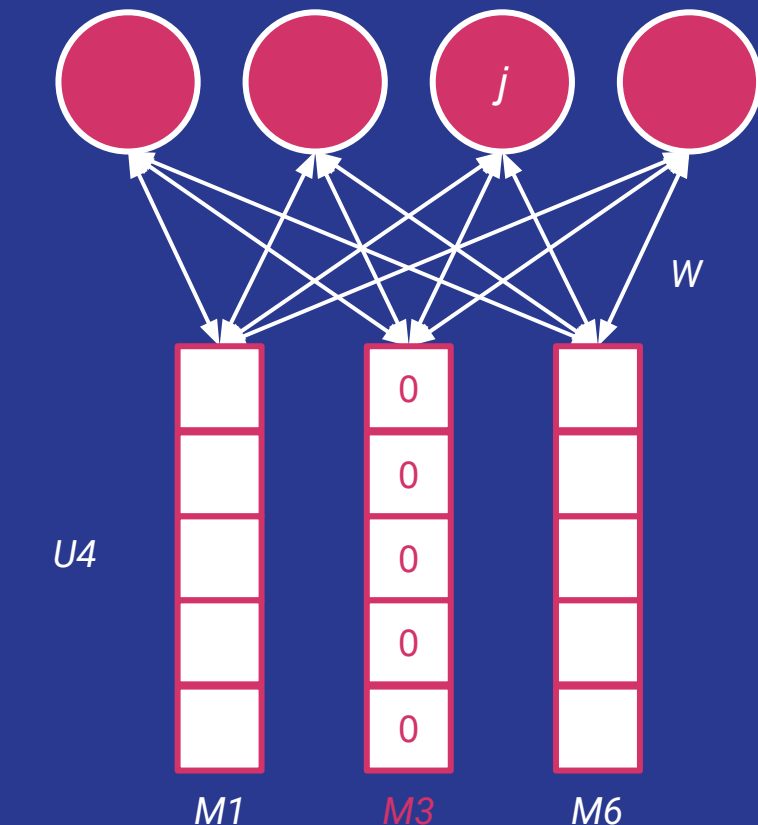
RBM Solution

Treat each user as a training case:

- A user is a vector of movie ratings.
- There is one visible unit per movie and its a 5-way softmax.
- The contrastive divergence learning rule for a softmax is the same as for a binary unit.
- There are ~ 100 hidden units.

One of the visible values is unknown:

- It needs to be filled in by the model.
- Softmax will yield the most probable rating (rating = $\text{argmax}(\mathbf{M3})$)



Training

- For each user, use an RBM that only has **visible units for the rated movies** (skip unrated - only for training!).
- Instead of one RBM for all users, we have a different RBM for every user.
 - All RBMs use the **same hidden units**.
 - The **weights** from each hidden unit to each movie **are shared** by all the users who rated that movie.
- Each user-specific RBM only gets **one training case!**
 - It works due to the weight-sharing.





Implementation

Initial Values

```
self.weights = torch.randn(num_v, num_h) * 0.01
self.v_bias = torch.ones(num_v) * 0.5
self.h_bias = torch.zeros(num_h)
self.momentum_coefficient = 0.5
self.k = 2
self.weight_decay = 1e-4
self.learning_rate = 1e-3
```

$$\mathbf{W} \sim \mathcal{N}(0, 0.01)$$

$$\mathbf{b}_v = 0.5 \text{ or } \log[p_i/(1-p_i)]$$

$$\mathbf{b}_h = 0$$

$$\alpha = 0.5$$

$$k \in [1, \infty]$$

$$\text{wd} = 0.0001$$

$$\varepsilon = 0.001$$

Hidden Vector

```
def sample_hidden(self, visible):  
    activations = torch.matmul(visible,  
                                self.weights) + self.bias_h  
    p = torch.sigmoid(activations)  
    return p
```

$$\mathbf{h}^{(0)} = \mathbb{E}[\mathbf{h} \mid \mathbf{v}] = \sigma(\mathbf{W}\mathbf{v}^{(0)} + \mathbf{b}_h)$$

Visible Vector

```
def sample_visible(self, hidden):  
    activations = torch.matmul(hidden,  
                                self.weights.t()) + self.bias_v  
    p = torch.sigmoid(activations)  
    return p
```

$$\mathbf{v}^{(1)} = \mathbb{E}[\mathbf{v} \mid \mathbf{h}] = \sigma(\mathbf{W}^T \mathbf{h}^{(0)} + \mathbf{b}_v)$$

Binarization

```
def binarize(self, samples):  
    p = torch.rand(self.num_h).to(self.device)  
    activated_units = (samples >= p).float()  
    return activated_units
```

$$p_i \sim U(0, 1)$$

unit = $\begin{cases} 1, & \text{if } s \geq p \\ 0, & \text{else} \end{cases}$

RBM Pass

```
def forward(self, input_data, train=True):  
    # Positive phase  
    pos_hidden_probabilities = self.sample_hidden(input_data)  
    hidden_activations = self.binarize(pos_hidden_probabilities)  
    positive_associations = torch.matmul(input_data.t(), hidden_activations)  
  
    # Negative phase  
    for step in range(self.k):  
        # Gibbs sampling  
        neg_visible_probabilities = self.sample_visible(hidden_activations)  
        neg_hidden_probabilities = self.sample_hidden(neg_visible_probabilities)  
        hidden_activations = self.binarize(neg_hidden_probabilities)  
    negative_associations = torch.matmul(neg_visible_probabilities.t(), hidden_activations)
```

Parameters Update


```
if train:
    # Update momentum
    self.weights_momentum *= self.momentum_coefficient
    self.weights_momentum += (positive_associations - negative_associations) * self.lr

    self.v_bias_momentum *= self.momentum_coefficient
    self.v_bias_momentum += torch.sum(input_data - neg_visible_probabilities, dim=0) * self.lr

    self.h_bias_momentum *= self.momentum_coefficient
    self.h_bias_momentum += torch.sum(pos_hidden_probabilities - neg_hidden_probabilities, dim=0) * self.lr

    # Update weights and biases
    batch_size = input_data.size(0)
    self.weights += self.weights_momentum / batch_size
    self.v_bias += self.v_bias_momentum / batch_size
    self.h_bias += self.h_bias_momentum / batch_size

    # L2 weight decay
    self.weights -= self.weights * self.weight_decay
```



Contrastive Divergence

```
def contrastive_divergence(self, input_data):  
    # Do k sampling steps and updates  
    input_data, negative_visible_probabilities =self.forward(input_data)  
  
    # Compute a reconstruction error  
    error = F.mse_loss(input_data, neg_visible_probabilities)  
  
    return error
```



Training

```
def train(model, train_loader, input_size, n_epochs=n_epochs):  
    for epoch in range(n_epochs):  
        epoch_error = 0.0  
        for batch, _ in train_loader:           # batch size = 64  
            batch = batch.view(1, input_size) # flatten input data  
            batch = batch.to(device)  
            batch_error = model.contrastive_divergence(batch)  
            epoch_error += batch_error  
  
        print(f'Epoch: {epoch + 1} | Error: {epoch_error:.4f}')
```

return model



Inference

```
images = next(iter(train_loader))[0]  
v, v_rec = model.forward(images.to(device).view(1, 784), train=False) # Flatten input data
```





Done!

Extra Reading

- Luis Serrano. Restricted Boltzmann Machines - A friendly introduction. https://www.youtube.com/watch?v=Fkw0_aAtwlw
- Geoffrey Hinton. A Practical Guide to Training Restricted Boltzmann Machines. <https://www.cs.toronto.edu/~hinton/absps/guideTR.pdf>
- Asja Fischer, Christian Igel. An Introduction to Restricted Boltzmann Machines. https://link.springer.com/chapter/10.1007/978-3-642-33275-3_2



The background is a solid pink color. In the top right corner, there is a decorative pattern of overlapping triangles in various shades of pink and magenta, creating a geometric, abstract design.

Thank you for
attention!